

```
import pandas as pd, numpy as np, matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

df = pd.read_csv('/content/housing.csv')

# Drop rows with any missing values
df.dropna(inplace=True)

print(df.head())
print(df.info())

# Separate features (X) and target (y)
X = df.drop(columns=['median_house_value'])
y = df['median_house_value']

# Apply one-hot encoding to 'ocean_proximity' in X
X = pd.get_dummies(X, columns=['ocean_proximity'], drop_first=True)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = LinearRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
print("MSE:", mean_squared_error(y_test, y_pred))

plt.figure()
plt.scatter(y_test, y_pred)
plt.xlabel("Actual")
plt.ylabel("Predicted")
plt.show()
```

```
longitude  latitude  housing_median_age  total_rooms  total_bedrooms \
0      -122.23     37.88            41.0        880.0         129.0
1      -122.22     37.86            21.0       7099.0        1106.0
2      -122.24     37.85            52.0       1467.0         190.0
3      -122.25     37.85            52.0       1274.0        235.0
4      -122.25     37.85            52.0       1627.0        280.0

population  households  median_income  median_house_value  ocean_proximity
0      322.0        126.0        8.3252      452600.0    NEAR BAY
1     2401.0       1138.0        8.3014      358500.0    NEAR BAY
2      496.0        177.0        7.2574      352100.0    NEAR BAY
3      558.0        219.0        5.6431      341300.0    NEAR BAY
4      565.0        250.0        2.8162      242200.0    NEAR BAY
```

Start coding or [generate](#) with AI.

In [1]: In [2]: Out[2]:

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	longitude	20433	non-null float64
1	latitude	20433	non-null float64
2	housing_median_age	20433	non-null float64
3	total_rooms	20433	non-null float64
4	total_bedrooms	20433	non-null float64
5	population	20433	non-null float64
6	households	20433	non-null float64
7	median_income	20433	non-null float64
8	median_house_value	20433	non-null float64
9	ocean_proximity	20433	non-null object

dtypes: float64(9), object(1)

memory usage: 1.7+ MB

None

MSE: 4802173538.60416

